

# Advanced Machine Learning: **Deep Generative Models**

## *Generative Adversarial Networks*

lecturer: Prof. Dr. Stephan Günnemann  
[www.daml.in.tum.de](http://www.daml.in.tum.de)

---

Summer Term 23


# Roadmap

---

- Deep Generative Models
  1. Introduction
  2. Normalizing Flows
  3. Variational Inference
  4. Variational Autoencoder
  - 5. Generative Adversarial Networks**
  6. Denoising Diffusion

# Learn a Generative Model with Fewer Assumptions

- The variational autoencoder assumes a latent variable structure and specific forms of the likelihood and the variational posterior
- Can we model the data better by having fewer assumptions?
- Idea: Transform given initial noise using a flexible function
  1. Draw an initial noise vector from the prior  $\mathbf{z} \sim p(\mathbf{z})$
  2. Deterministically transform  $\mathbf{z}$  into the target data/output space  $\mathbf{x} = f_{\theta}(\mathbf{z})$



- In contrast to Normalizing Flows,  $f_{\theta}(\mathbf{z})$  can be any function

# Resulting Distribution

1. Draw an initial noise vector from the prior  $\mathbf{z} \sim p(\mathbf{z})$
2. Deterministically transform  $\mathbf{z}$  into the target data/output space  $\mathbf{x} = f_{\theta}(\mathbf{z})$

- This process defines a valid density on the output space

$$p_{\theta}(\mathbf{x}) = \frac{d}{dx_1} \cdots \frac{d}{dx_d} \int_{\{f_{\theta}(\mathbf{z}) \leq \mathbf{x}\}} p(\mathbf{z}) d\mathbf{z}$$

- Unfortunately  $p_{\theta}(\mathbf{x})$  is intractable to compute
  - The integration region  $\{f_{\theta}(\mathbf{z}) \leq \mathbf{x}\}$  itself may be very hard to determine
  - The integral itself may then be computationally demanding
  - Taking  $d$  partial derivatives in a high-dimensional space could be challenging

# Learn a Generative Model with Fewer Assumptions

1. Draw an initial noise vector from the prior  $\mathbf{z} \sim p(\mathbf{z})$
  2. Deterministically transform  $\mathbf{z}$  into the target data/output space  $\mathbf{x} = f_{\theta}(\mathbf{z})$
- Sampling is easy; density evaluation is hard!
- What can we do if we can only sample but we cannot evaluate the density?
    - Without a (tractable) **likelihood function**, many widely-used tools for inference and parameter learning become **unavailable** (e.g. the previously seen variational inference)

# Forward Pointer: Generative Adversarial Networks

- GANs have been very popular for learning deep generative models
- Informally, the main idea is:

- Two competing neural network models

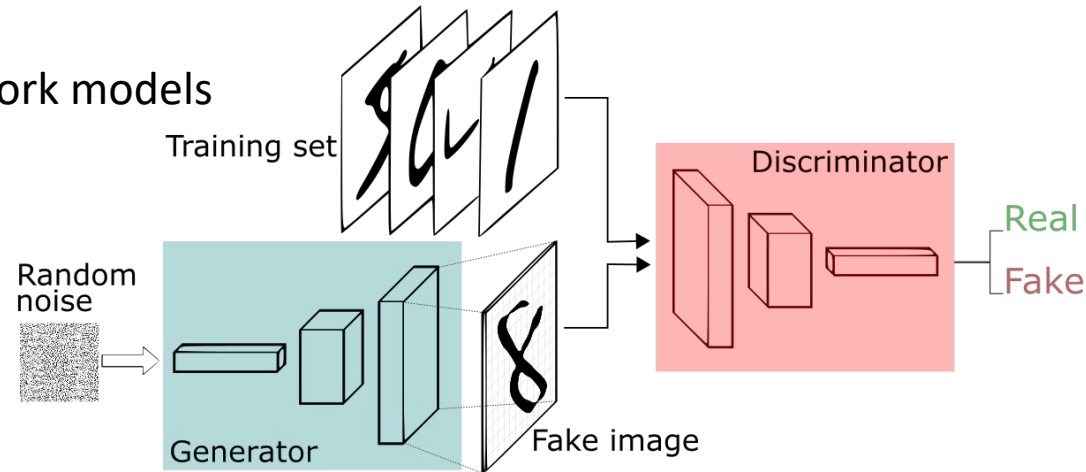
- Generator: takes noise as input and generates ("fake") samples

- Discriminator: receives samples from both generator and training data

and has to distinguish between the two → classify input as "real" or "fake"

- Goal: Train the generator in such a way that the discriminator can **not** distinguish between real and "fake" samples

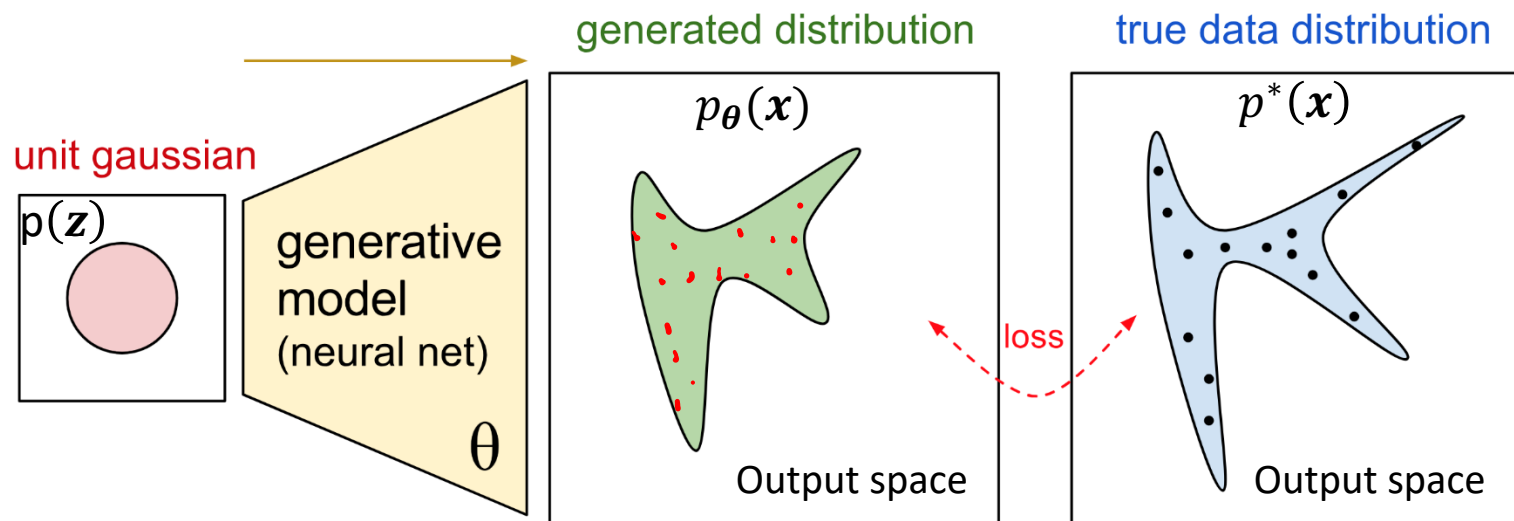
- In this case, the generator generates realistic examples




In the following, you learn the actual (and more general) concept behind it

# Likelihood-free inference

- What can we do if we can **easily draw samples** from the model but we cannot evaluate the density?
- Idea: We can use **any** method that compares two **sets of samples**.
  - This process is called ***density estimation by comparison***.
- We **test** the **hypothesis** that the true data distribution  $p^*(x)$  and the model distribution  $p_\theta(x)$  are **equal**




# Density Estimation by Comparison




**Density Difference**


$$r_\phi = p^* - p_\theta$$



*Max Mean  
Discrepancy*




*Moment  
Matching*




*Bregman  
Divergence*

**Density Ratio**

$$r_\phi = \frac{p^*}{p_\theta}$$



*Class Probability  
Estimation*



*f-Divergence*

Figure: [Mohamed and Lakshminarayanan, 2016]



# Density Difference Approaches (I)

- Moment Matching
  - Use the fact that  $p^*$  and  $p_\theta$  are identical iff the expectations of any test statistic  $s(\mathbf{x})$  are identical
  - Thus, minimize the gap  $\min_{\theta} (\mathbb{E}_{p^*(\mathbf{x})}[s(\mathbf{x})] - \mathbb{E}_{p(\mathbf{z})}[s(f_\theta(\mathbf{z}))])^2$
  - Approximate expectation by Monte-Carlo samples

# Density Difference Approaches (II)

---

- Max Mean Discrepancy
  - When the statistics  $s(\mathbf{x})$  are defined within a reproducing kernel Hilbert space, we obtain kernel-based forms of these objectives
  - These kernels are highly flexible and allow easy handling of data such as images, strings and graphs

# Ratio-Based Approaches (I)

- Density ratio  $r^*(\mathbf{x}) = p^*(\mathbf{x}) / p_\theta(\mathbf{x})$ 
  - In the best case always 1, i.e. the two distributions are indistinguishable
  - However, we cannot compute ratio in closed form/easily
- Idea: Approximate the true density ratio  $r^*(\mathbf{x})$  by  $r_\phi(\mathbf{x})$ 
  - Finding the approximation  $r_\phi(\mathbf{x})$  often means solving again a learning problem
- Thus, we get the following general principle for learning
  - **Optimize Ratio loss:**  
approximate the true density ratio  $r^*(\mathbf{x})$  (i.e. learning  $\phi$ )
  - **Optimize Generative loss:**  
drive the density ratio towards 1 (i.e. learning  $\theta$ )
  - Essentially a bi-level optimization problem, which is usually just solved alternatingly

# Ratio-Based Approaches (II)

- The following three methods are an instantiation of this principle:
- Class Probability Estimation
  - Estimate the density ratio via a classifier which can distinguish between the observed data and data generated from the model (principle used in, e.g., **Generative Adversarial Networks (GANs)**)
- Divergence Minimization
  - Minimize an  $f$ -divergence between the true density and the generator density
  - For specific choices of the divergence function this can be equivalent to the class probability estimation case or KL divergence
- Ratio matching
  - Directly minimize the error between the true density ratio and an estimate of it:
$$\mathcal{L} = \mathbb{E}_{p_\theta} \left[ \left( r_\phi(\mathbf{x}) - r^*(\mathbf{x}) \right)^2 \right]$$
  - Can be generalized beyond the squared error using the Bregman divergence

# Learning via Class Probability Estimation (I)

- Let  $Y$  denote a random variable which assigns label  $Y = 1$  to samples from the true data distribution;  $Y = 0$  to those from the generator distribution
- Then  $p(\mathbf{x} \mid Y = 1) = p^*(\mathbf{x})$  and  $p(\mathbf{x} \mid Y = 0) = p_\theta(\mathbf{x})$

- Denote  $P(Y = 1) = \pi$ . From Bayes we have:

$$r^*(\mathbf{x}) = \frac{p^*(\mathbf{x})}{p_\theta(\mathbf{x})} = \frac{p(Y = 1 \mid \mathbf{x})}{p(Y = 0 \mid \mathbf{x})} \frac{1 - \pi}{\pi}$$

- Apparently, density ratio estimation is equal to **class-probability estimation**
  - Simply speaking: we can consider a classifier for  $\mathbf{x}$  (predicting labels  $Y=0$  or  $1$ )
- Specify a scoring function or a **discriminator**  $D_\phi(\mathbf{x}) = p(Y = 1 \mid \mathbf{x})$
- e.g. logistic regression or a neural network

# Learning via Class Probability Estimation (II)

- Specify a scoring function or a **discriminator**  $D_\phi(\mathbf{x}) = p(Y = 1 | \mathbf{x})$ 
  - that tells you whether a sample  $\mathbf{x}$  is real or from the generator (“fake”)
  - e.g. logistic regression or a neural network
- For learning  $D_\phi(\mathbf{x})$ , we need a loss function, e.g., the cross-entropy loss:

$$\begin{aligned}\mathcal{L}_{\theta, \phi} &= \mathbb{E}_{(x, y) \sim p(x, y)} \left[ -y \log[D_\phi(\mathbf{x})] - (1 - y) \log[1 - D_\phi(\mathbf{x})] \right] \\ &= \pi \mathbb{E}_{p^*(\mathbf{x})} \left[ -\log D_\phi(\mathbf{x}) \right] + (1 - \pi) \mathbb{E}_{p(\mathbf{z})} \left[ -\log[1 - D_\phi(f_\theta(\mathbf{z}))] \right]\end{aligned}$$

# Learning via Class Probability Estimation (III)

1. Solving 
$$\phi^*(\theta) = \underset{\phi}{\operatorname{argmin}} \mathcal{L}_{\theta, \phi}$$
 leads to the "best" discriminator for a given generative model ( $\theta$ )
  - That is, we will approximate  $r^*(\mathbf{x}) = \frac{p^*(\mathbf{x})}{p_{\theta}(\mathbf{x})} = \frac{p(Y=1 | \mathbf{x})}{p(Y=0 | \mathbf{x})} \approx \frac{D_{\phi^*(\theta)}(\mathbf{x})}{1 - D_{\phi^*(\theta)}(\mathbf{x})}$ 
    - here w.l.o.g. we set  $\pi = 0.5$
2. We aim to drive the density ratio  $r^*(\mathbf{x})$  towards 1
  - aim:  $p(Y = 1 | \mathbf{x}) = p(Y = 0 | \mathbf{x})$
  - That is, find generative model ( $\theta$ ) such that (even) the "best" discriminator cannot distinguish the classes
  - Solving 
$$\theta^* = \underset{\theta}{\operatorname{argmax}} \mathcal{L}_{\theta, \phi^*(\theta)}$$

# Learning via Class Probability Estimation (IV)

Attention: Note  
the flipped sign  
in the loss  
formulation!

$$-\mathcal{L}_{\theta,\phi}$$

- Generator and discriminator play a **minimax game**:

$$\min_{\theta} \max_{\phi} \pi \mathbb{E}_{p^*(\mathbf{x})} [\log D_{\phi}(\mathbf{x})] + (1 - \pi) \mathbb{E}_{p(\mathbf{z})} [\log [1 - D_{\phi}(f_{\theta}(\mathbf{z}))]]$$

- Discriminator: aims to distinguish between (samples from)  $p^*(\mathbf{x})$  and  $p_{\theta}(\mathbf{x})$

➤ Maximization

// minimization of cross-entropy

- Generator: aims to generate samples that are indistinguishable

➤ Minimization

// maximization of (lowest) cross-entropy

- This is a bilevel optimization problem



# Learning via Class Probability Estimation (V)

- This bilevel problem is typically tackled via alternating optimization
  - usually does not lead to the optimal solution
  - depending on the setting, difficult to train (instable)

- **Ratio loss** (discriminator loss) optimization:

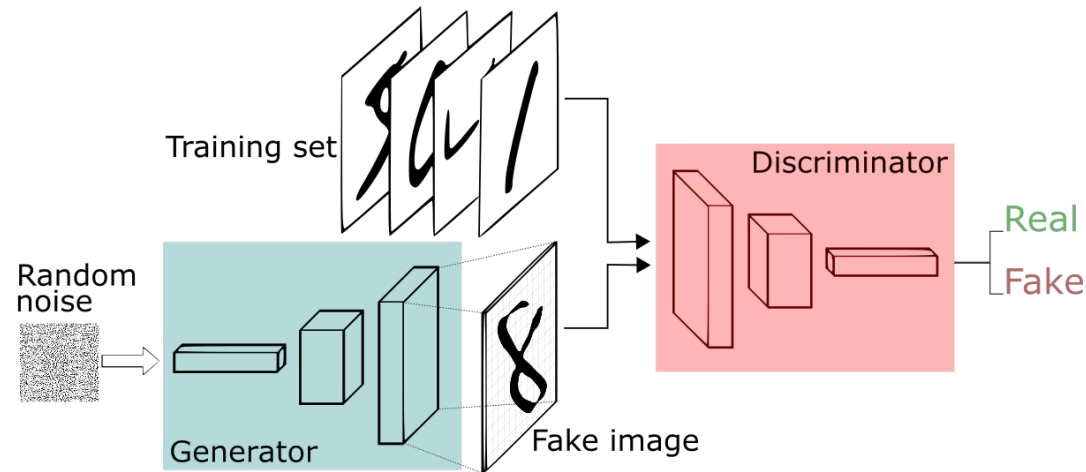
$$\min_{\phi} \pi \mathbb{E}_{p^*(\mathbf{x})} [-\log D_{\phi}(\mathbf{x})] + (1 - \pi) \mathbb{E}_{p(\mathbf{z})} [-\log [1 - D_{\phi}(f_{\theta}(\mathbf{z}))]]$$

- **Generative loss** optimization:

$$\min_{\theta} \mathbb{E}_{p(\mathbf{z})} [\log [1 - D_{\phi}(f_{\theta}(\mathbf{z}))]]$$

# GANs – Generative Adversarial Networks

- Main idea
  - Two competing neural network models
  - Generator: takes noise as input and generates ("fake") samples
  - Discriminator: receives samples from both generator and training data and has to distinguish between the two
  - They play a minimax game against each other
  - Competition drives the generated samples to be indistinguishable from real



- In short: GAN approach =  
Learning of a generator via class probability estimation where the generator and the discriminator are neural networks

# GANs – Examples

- Synthetically generated images



Figure 5:  $1024 \times 1024$  images generated using the CELEBA-HQ dataset. See Appendix F for a larger set of results, and the accompanying video for latent space interpolations.

From "Progressive Growing of GANs for Improved Quality, Stability, and Variation", Tero Karras et.al.

# Questions – GANs

---

1. What can we say about the ratio  $r^*(\mathbf{x}) = p^*(\mathbf{x}) / p_\theta(\mathbf{x})$  when:
  - a) The generator and the discriminator are optimal
  - b) The generator only is optimal
  - c) The discriminator only is optimal
  
2. For a given data  $\mathbf{x}$ , what is the probability for the discriminator to be correct when:
  - a) The generator and the discriminator are optimal
  - b) The generator only is optimal
  - c) The discriminator only is optimal

# References

---

- [1] Mohamed, Shakir, and Balaji Lakshminarayanan. “Learning in Implicit Generative Models.”: <https://arxiv.org/pdf/1610.03483.pdf>
- [2] Andrew Davison blog :  
<https://casmls.github.io/general/2017/05/24/ligm.html>

# External Sources

---

## Web tutorial:

- Andrew Davison blog : <https://casmls.github.io/general/2017/05/24/ligm.html>  
Eric Jang blog : <https://blog.evjang.com/2018/01/nf1.html>

## Survey papers:

- Mohamed, Shakir, and Balaji Lakshminarayanan. “Learning in Implicit Generative Models.”: <https://arxiv.org/pdf/1610.03483.pdf>