

# Advanced Machine Learning: **Deep Generative Models**

## *Introduction*

Lecturer: Prof. Dr. Stephan Günnemann

[www.daml.in.tum.de](http://www.daml.in.tum.de)

---

Summer Term 2023

Data Analytics and  
Machine Learning



# Roadmap

---

- Introduction
  1. **What will you learn in this lecture?**
  2. Organizational aspects + project tasks

# What is this course about?

---

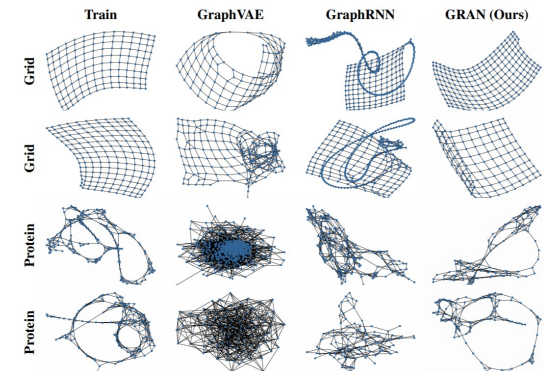
- In short: A continuation of our intro ML lecture (IN2064) now focusing on advanced learning principles and deep generative models
- Focus on algorithms and general principles, not limited to a single domain
- Project tasks will give you hands-on experience
- At the end you should also be able to extend existing techniques and adapt them to different domains and applications

# How do you learn complex distributions?

- How can we learn probability distributions over complex real-world data such as images, graphs, and audio signals?





[Khan, 2019]



[Liao+, 2019]


- Distributions in high dimensions defy our intuition



- How do we design flexible and efficient models for these settings?

# Generative Models

- Deterministic Generative Models
  - Image = Renderer(object=cube, color=red, size=, position=, ...)
  - Image = Renderer(object=cylinder, color=blue, size=, position=, ...)
- Statistical Generative Models



+

- Model family
- Loss function
- Optimization algorithm
- ....

learning

=

 $p(x)$ 

Data

Prior Knowledge

Probability Distribution

# Desiderata for Statistical Generative Models

## 1. Efficient Sampling

- Should be easy to sample a new instance  $\mathbf{x}_{new} \sim p(\mathbf{x})$
- Sampled/generated instances  $\mathbf{x}_{new}$  should be similar to the training data

## 2. Efficient Likelihood Evaluation

- Should be easy to evaluate  $p(\mathbf{x})$  for any instance  $\mathbf{x}$ , e.g.  $p(\img alt="A golden retriever puppy" data-bbox="718 458 842 568"/>$

## 3. (Optionally) Extract Features

- For any instance  $\mathbf{x}$  extract latent features/representations
- Capture/summarize the important aspects of the instance/image

# Some Applications of Generative Models

- Image generation

- <https://www.thispersondoesnotexist.com/>
- <https://www.youtube.com/watch?v=p5U4NgVGAwg>


- 3D graphics & fluid dynamics

- <https://www.youtube.com/watch?v=i6JwXYypZ3Y>

- Speech & music synthesis

- <https://deepmind.com/blog/article/wavenet-generative-model-raw-audio>

- Drug discovery



[Lim+, 2019]




[Achlioptas+, 2017]



[Karras+, 2018]

# Continuous Distributions over High-dimensional Data

- “Classic” probability distributions (e.g. multivariate normal) do not capture the complexity of real-world datasets
  - Real distributions are multi-modal, asymmetric




- Can we use mixture models to capture this behavior?



# Mixture models

- In theory, a mixture with enough components can represent any density
  - How many is “enough”?



[Vergari+, 2019]

- Even for simple 2D densities we need hundreds of mixture components!
  - The situation gets (exponentially) worse as we increase the dimensionality

# Discrete Distributions over High-dimensional Data

- What about discrete distributions?
- Suppose  $x_1, x_2, x_3$  are binary variables
  - $p(x_1, x_2, x_3)$  can be specified with  $2^3 - 1 = 7$  parameters
  - $p(0,0,0), p(0,0,1), \dots p(1, 1, 0), \cancel{p(1,1,1)}$
- For an image with  $N$  black or white pixels need to specify  $2^N - 1$  values
  - The number of parameters grows exponentially with dimension

# Challenges of High-dimensional Data

---

- “Classic” distributions
  - Do not capture the complexity of the data
- (Finite) mixture models
  - Require ridiculous amounts of parameters to specify even simple densities
  - Do not work in higher dimensions
- For discrete distributions – combinatorial explosion
  
- In this lecture, you will learn how to design flexible and efficient generative models for high-dimensional data using deep learning techniques

# Contents of this Course

---

- Deep Generative Models
  - Normalizing flows
  - Variational inference / Variational Autoencoder
  - Generative Adversarial Networks
  - Denoising Diffusion

# Roadmap

---

- Chapter: Introduction
  1. What will you learn in this lecture?
  2. **Organizational aspects + project tasks**

# Course Organization

---

- Lecturer
  - Prof. Dr. Stephan Günnemann
- Teaching assistants
  - Dominik Fuchsgruber, Marten Lienen, David Lüdke
- 3 ECTS
- Language: English
- Ungraded exercise sheets
- Graded project tasks
- Final exam + repeat exam
- Project tasks can grant a bonus of up to 0.3

# Schedule

---

- In-person lectures and exercises
  - Lecture Thursday 09:00 – 11:00
  - Exercise roughly every third Thursday 09:00 – 11:00
- Practice material and exercises uploaded to Moodle
  - Ungraded exercise sheets every third week
  - Project tasks every third week

# Preliminary Timetable

Week		Topic	Project
1	20.04.2023	<b>Orga / Normalizing Flow</b>	
2	27.04.2023	Exercise	Normalizing Flows
3	04.05.2023	<b>Variational Inference</b>	
4	11.05.2023	<b>Variational Inference</b>	
No lecture and Exercise			
5	25.05.2023	Exercise	
6	01.06.2023	<b>VAE</b>	
No lecture and Exercise			
7	15.06.2023	<b>VAE/GANs</b>	
8	22.06.2023	<b>GANs</b>	VAE's
9	29.06.2023	Exercise	
10	06.07.2023	<b>Diffusion</b>	
11	13.07.2023	<b>Diffusion</b>	Diffusion
12	20.07.2023	Exercise	



# Prerequisites

---

- The course is designed for Master students of Computer Science (and specializations such as Data Engineering and Analytics, Games Engineering, etc.)
  
- **This course can not be taken by students that passed MLGS in previous years**
  
- Prerequisites:
  - Knowledge about the standard Machine Learning concepts (i.e. content of our lecture IN2064)
    - We assume the basic concepts are clear; no repetition!
    - We strongly recommend that you attend IN2064 first before taking this class
  - Knowledge about:
    - Algorithms and Data structures
    - Programming
    - Mathematics: Linear Algebra, Statistics, Optimization

# Course Material + Announcements

---

- All course materials (slides, exercises) will be uploaded to Moodle
  - Video recordings of previous years lectures are accessible via link on Moodle
- Project submission on Artemis
- Use Piazza to ask questions! (please avoid sending e-mails)

<https://piazza.com/tum.de/summer2023/cit4230003>

Access Code: **dgm2023**

- Please read the [guidelines](#) for using Piazza

# Exercises and Project Tasks

---

- Exercise sheets
  - Exam preparation
  - Solutions in the tutorials
  - Due to the high number of registrations, we are unable to provide corrections to your solutions
  
- Project tasks
  - Get hands-on experience with advanced machine learning methods
  - Improve your final grade! (details later)

# Project Format

---

- Format of programming tasks
  - Tasks will be published via Artemis
  - [artemis.in.tum.de](https://artemis.in.tum.de)
  
- How to solve programming tasks?
  - Clone template repository from Artemis exercise
  - Solve tasks described in the repository
  - Push repository with filled-in solutions
  
- Bonus regulations
  - 10 points for each programming sheet
  - A Bonus of 0.1, 0.2 and 0.3 grade points will be granted upon correct completion of 25%, 50% and 75% of all project points, respectively

# Project topics

---

- Three project tasks on the following topics
  - Normalizing flows
  - Variational Autoencoder
  - Denoising Diffusion
- The specific tasks and all details will be described in the corresponding descriptions on Artemis

# Exam & Grading Scheme

- Written final exam: 90 minutes
  - Date will be announced via TUMonline
  - We currently plan with an on-site exam
  - One handwritten two-sided A4 sheet with notes

```
def final_grade(exam_grade, project_grade):  
    if exam_grade > 4.0:  
        return exam_grade  
    else:  
        return max(1.0, exam_grade - bonus)
```

→ The project is voluntary and can only improve the final grade. The project bonus applies only if you passed, and you cannot improve beyond 1.0

# Our Group's Focus

## Reliable Machine Learning for **Non-Independent** Data



### Robustness/Uncertainty

- Data corruptions, adversaries
  - Certificates

### Non-independent data

- Temporal/sequence data
  - Graph data

- Interested? We offer:
  - Bachelor/Master theses, Guided Research projects, HiWi positions
- More details on specific topics closer to the end of the semester

# References

---

Figures taken from

- Goodfellow et al. 2014, <https://arxiv.org/abs/1412.6572>
- Akbik et al. 2018, <https://research.zalando.com/welcome/mission/research-projects/flair-nlp/>
- Khan 2019, <https://heartbeat.fritz.ai/stylegans-use-machine-learning-to-generate-and-customize-realistic-images-c943388dc672>
- Liao et al. 2019, <https://arxiv.org/abs/1910.00760>
- Lim et al. 2019, Scaffold-based molecular design with a graph generative model
- Achlioptas et al. 2017, <https://arxiv.org/abs/1707.02392>
- Karras et al. 2019, <https://github.com/NVlabs/stylegan>
- Vegari et al. 2019, <https://web.cs.ucla.edu/~guyvdb/slides/TPMTutorialUAI19.pdf>