


Introduction to Deep Learning (I2DL)


Exercise 5: Neural Networks and CIFAR10 Classification

Today's Outline

- Neural Networks
 - Mathematical Motivation
 - Modularization
- Exercise 5
 - Implementation Loop
 - CIFAR10 Classification



Our Goal




\$?

Universal Approximation Theorem

Theorem (1989, colloquial)

For any continuous function f on a compact set K , there exists a one layer neural network, having only a single hidden layer + sigmoid, which uniformly approximates f to within an arbitrary $\varepsilon > 0$ on K .



Universal Approximation Theorem (Optional)

Readable proof:


https://mcneela.github.io/machine_learning/2017/03/21/Universal-Approximation-Theorem.html

(Background: Functional Analysis, Math Major 3rd semester)

Visual proof:

<http://neuralnetworksanddeeplearning.com/chap4.html>


A word of warning...




Source: <http://blog.datumbox.com/wp-content/uploads/2013/10/gradient-descent.png>

Shallow vs. Deep

- Shallow
(1 hidden layer)



- Deep
(>1 hidden layer)



Obvious Questions

- Q: Do we even need deep networks?
A: Yes. Multiple layers allow for more representation power given a fixed computational budget in comparison to a single layer
- Q: So we just build 100 layer deep networks?
A: Not trivially ;-)
Constraints: Memory, vanishing gradients, ...

Exercise 4: Simple Classification Net

```
class Classifier(Network):
    """
    Classifier of the form  $y = \text{sigmoid}(X * W)$ 
    """

    def __init__(self, num_features=2):
        super(Classifier, self).__init__("classifier")

        self.num_features = num_features
        self.W = None

    def initialize_weights(self, weights=None):
        """
        Initialize the weight matrix W

        :param weights: optional weights for initialization
        """
        if weights is not None:
            assert weights.shape == (self.num_features + 1, 1), \
                "weights for initialization are not in the correct shape"
            self.W = weights
        else:
            self.W = 0.001 * np.random.randn(self.num_features + 1, 1)
```

```
    def forward(self, X):
        """
        Performs the forward pass of the model.

        :param X: N x D array of training data. Each row is a D-dimensional point.
        :return: Predicted labels for the data in X, shape N x 1
        """
        # 1-dimensional array of length N with classification scores.
        # =====
        assert self.W is not None, "weight matrix W is not initialized"
        # add a column of 1s to the data for the bias term
        batch_size, _ = X.shape
        X = np.concatenate((X, np.ones((batch_size, 1))), axis=1)
        # save the samples for the backward pass
        self.cache = X
        # output variable
        y = None
        # =====
        # TODO: Implement the forward pass and return the output of the model. Note
        # that you need to implement the function self.sigmoid() for that
        # =====
        y = X.dot(self.W)
        y = self.sigmoid(y)

        # =====
        #                               END OF YOUR CODE
        # =====
```

Modularization

Chain Rule:

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial d} \cdot \frac{\partial d}{\partial y}$$




```
class Sigmoid:
    def __init__(self):
        pass

    def forward(self, x):
        """
        :param x: Inputs, of any shape


        :return out: Output, of the same shape as x
        :return cache: Cache, for backward computation, of the same shape as x
        """

    def backward(self, dout, cache):
        """
        :return: dx: the gradient w.r.t. input X, of the same shape as X
        """
```


Exercise 3: Dataset




Exercise 4: Binary Classification



Exercise 5: Neural Networks and CIFAR10 Classification



Exercise 6: Neural Networks and Hyperparameter Tuning



Summary

- Monday 22.11: Watch Lecture 6
 - Training NNs
- Wednesday 24.11 15:59: Submit exercise 5
- Thursday 25.11: Tutorial 6
 - Hyper-parameter Tuning

See you next week

