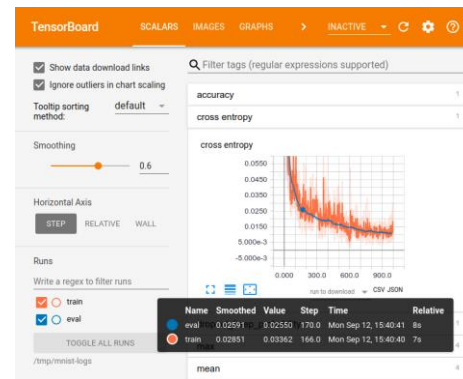


Introduction to Deep Learning (I2DL)

Exercise 7: Pytorch

Overview

- Exercise 6
 - Leaderboard
 - Case Study
- Deep Learning Frameworks
 - Static vs Dynamic
- Exercise 7
 - Pytorch, Tensorboard, Pytorch Lightning
- Outlook and Exam



Exercise 6 – Leaderboard

Exercise 1	Exercise 3	Exercise 4	Exercise 5	Exercise 6	Exercise 7	Exercise 8	Exercise 9	Exercise 10	Exercise 11
#	User	Score							
1	u1051	62.88							
2	u1048	59.08							
3	u1785	57.45							
4	u1120	56.75							
5	u0111	56.59							
6	u1005	56.32							

CIFAR10 - Leaderboard

RANK	MODEL	PERCENTAGE CORRECT [↑]	PERCENTAGE ERROR	FLOPS	PARAMS	ECE	EXTRA TRAINING DATA	PAPER	CODE	RESULT	YEAR
1	EffNet-L2 (SAM)	99.70	0.30				✓	Sharpness-Aware Minimization for Efficiently Improving Generalization	🔗	📄	2020
2	ViT-H/14	99.5	0.5		632M		✓	An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale	🔗	📄	2020
3	ViT-L/16	99.42	0.58		307M		✓	An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale	🔗	📄	2020
4	BiT-L (ResNet)	99.37	0.63				✓	Big Transfer (BiT): General Visual Representation Learning	🔗	📄	2019
5	LaNet	99.03	0.97		45		×	Sample-Efficient Neural Architecture Search by Learning Action Space for Monte Carlo Tree Search	🔗	📄	2019

Case Study: Power of Transformations

In [5]: `import random`

```
class SaltPepper:
    def __init__(self, grains=64, prob=0.5):
        self.grains = grains
        self.prob = prob

    def __call__(self, image):
        if random.uniform(0,1) < self.prob:
            rows = np.random.randint(0, image.shape[0], size=self.grains)
            cols = np.random.randint(0, image.shape[1], size=self.grains)

            g = self.grains // 2
            image[rows[g], cols[g]] = 255
            image[rows[g], cols[g]] = 0


        return image
```

In [6]: `#Retrieve an image from the dataset and flip it`

```
image = dataset[1890]['image']
transform = SaltPepper(64,1.0)
image_flipped = transform(image)

#Show the two images
plt.figure(figsize = (2,2))
plt.subplot(1, 2, 1)
plt.imshow(image.astype('uint8'))
plt.axis('off')
plt.subplot(1, 2, 2)
plt.imshow(image_flipped.astype('uint8'))
plt.axis('off')
plt.title("Left: Original Image, Right: Flipped image")
plt.show()
```

Left: Original Image, Right: Flipped image



```
In [7]: class RandomRoll:
        def __init__(self, dx=(-8,8), dy=(-8,8), prob=0.5):
            self.dx = dx
            self.dy = dy
            self.prob = prob

        def __call__(self, image):
            if random.uniform(0,1) < self.prob:
                x = random.randint(self.dx[0], self.dx[1])
                y = random.randint(self.dy[0], self.dy[1])
                return np.roll(image, (x, y), axis=(0,1))
            return image
```

In [8]: `#Retrieve an image from the dataset and flip it`

```
image = dataset[1]['image']
transform = RandomRoll((-8,8), (-8,8), 1.0)
image_flipped = transform(image)

#Show the two images
plt.figure(figsize = (2,2))
plt.subplot(1, 2, 1)
plt.imshow(image.astype('uint8'))
plt.axis('off')
plt.subplot(1, 2, 2)
plt.imshow(image_flipped.astype('uint8'))
plt.axis('off')
plt.title("Left: Original Image, Right: Flipped image")
plt.show()
```

Left: Original Image, Right: Flipped image



Case Study: Power of Transformations

```
In [9]: class RandomSpeckle:
def __init__(self, std=0.1, prob=0.5):
    self.std = std
    self.prob = prob

def __call__(self, image):
    if random.uniform(0,1) < self.prob:
        row,col,ch = image.shape
        gauss = np.random.randn(row,col,ch)
        gauss = gauss.reshape(row,col,ch)
        noisy = image + image * random.uniform(0.0, self.std) * gauss
        return noisy
    return image
```

```
In [10]: #Retrieve an image from the dataset and flip it
image = dataset[1000]['image']
transform = RandomSpeckle(prob=1.0)
image_flipped = transform(image)

#Show the two images
plt.figure(figsize = (2,2))
plt.subplot(1, 2, 1)
plt.imshow(image.astype('uint8'))
plt.axis('off')
plt.subplot(1, 2, 2)
plt.imshow(image_flipped.astype('uint8'))
plt.axis('off')
plt.title("Left: Original Image, Right: Flipped image")
plt.show()
```

Left: Original Image, Right: Flipped image



Final Training with parameters:

- num_layers
- hidden_size
- Regularization
- Learning rate

```
In [18]: trafo = ComposeTransform([
    rescale_transform,
    normalize_transform,
    RandomHorizontalFlip(),
    RandomSpeckle(prob=0.4),
    SaltPepper(prob=0.4),
    RandomRoll(prob=0.4),
    flatten_transform])

traf_dataset = ImageFolderDataset(
    mode='train',
    root=cifar_root,
    download_url=download_url,
    transform=trafo
)

data = DataLoader(
    dataset=traf_dataset,
    batch_size=512,
    shuffle=True,
    drop_last=True
)

In [19]: best_model = ClassificationNet(activation=Relu(), num_layer=3, hidden_size=258, reg=6.357358814407898e-06)
solver = Solver(best_model, data, dataloaders['val'], learning_rate=5.535943214549563e-05)
solver.train(epochs=500, patience=21)
```

Limiting Factors

- Computational Power
- Specialized Architectures
 - CNN
- Proper Initialization

RANK	MODEL	PERCENTAGE CORRECT [†]	PERCENTAGE ERROR	FLOPS	PARAMS	ECE	EXTRA TRAINING DATA	PAPER	CODE	RESULT	YEAR
1	EffNet-L2 (SAM)	99.70	0.30				✓	Sharpness-Aware Minimization for Efficiently Improving Generalization	🔗	📄	2020
2	ViT-H/14	99.5	0.5		632M		✓	An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale	🔗	📄	2020
3	ViT-L/16	99.42	0.58		307M		✓	An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale	🔗	📄	2020
4	BiT-L (ResNet)	99.37	0.63				✓	Big Transfer (BiT): General Visual Representation Learning	🔗	📄	2019
5	LaNet	99.03	0.97		45		×	Sample-Efficient Neural Architecture Search by Learning Action Space for Monte Carlo Tree Search	🔗	📄	2019

Deep Learning Frameworks

The two big ones

- Tensorflow - Google
 - As well as Keras
- Pytorch - Facebook




Other examples

- CNTK - Microsoft
- Mxnet - Apache
- Jax - Google

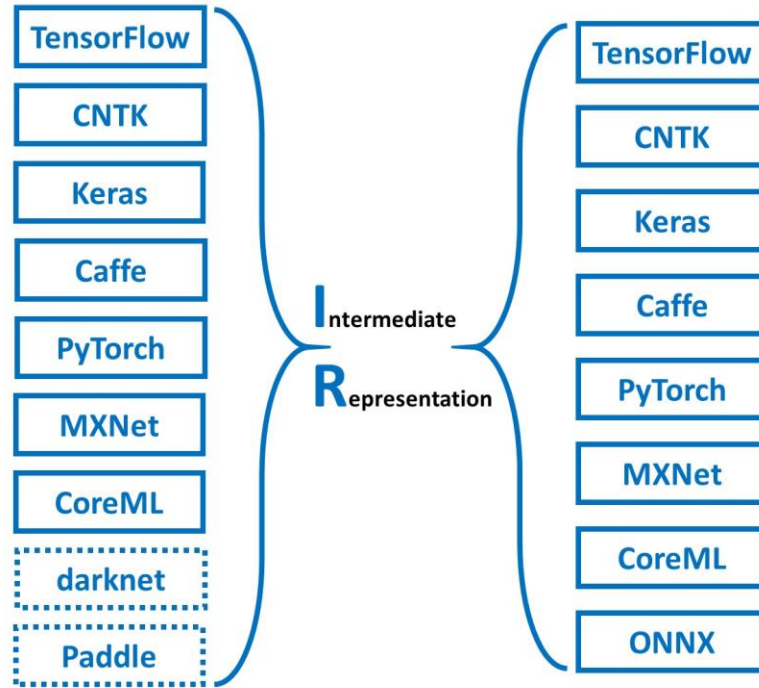


Handling of Computation Graph



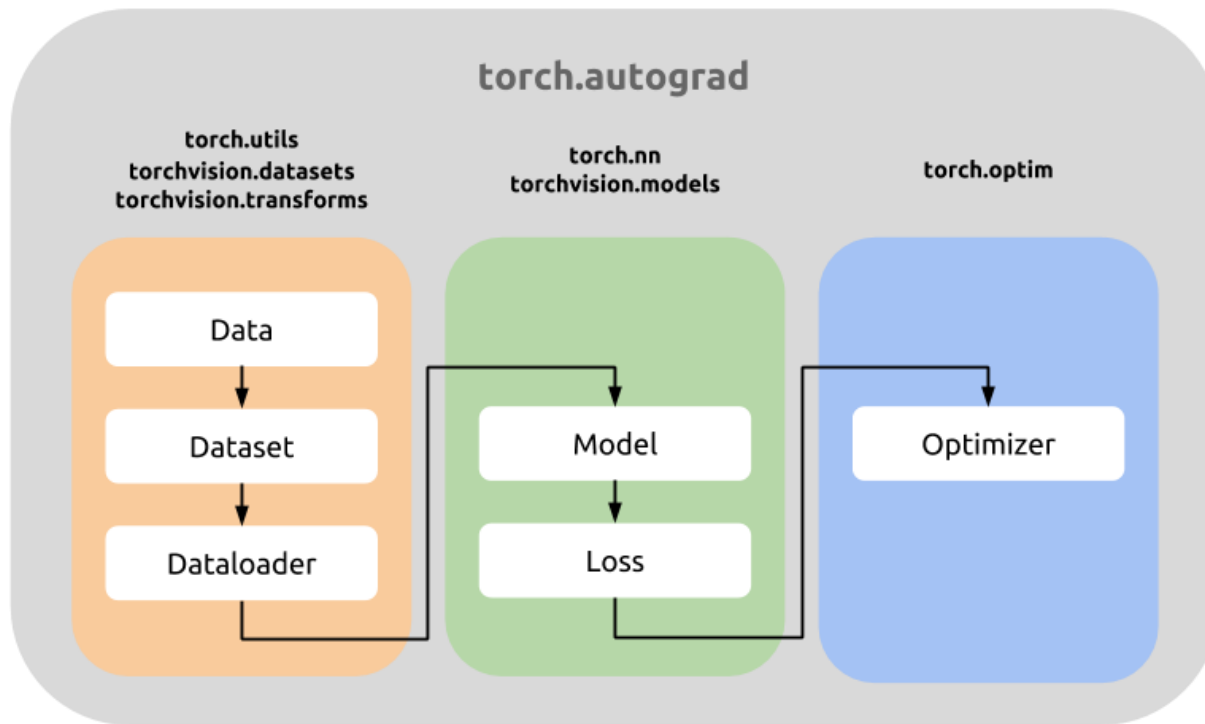
	Tensorflow	Pytorch
Graph Creation	Static/Eager	Dynamic/On Runtime
Similar to	C	Python

Framework Conversion



See: <https://github.com/microsoft/MMdnn>

Pytorch: Overview



Easy Device Assignment

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

print(f"Original device: {x.device}") # "cpu", integer


tensor = x.to(device)
print(f"Current device: {x.device}") # "cpu" or "cuda", double
```

```
cpu
Original device: cpu
Current device: cpu
```

Datasets: Torchvision

- Torchvision
 - `torchvision.datasets` contains many datasets, such as `ImageNet`, `FashionMNIST`, etc.

```
fashion_mnist_dataset = torchvision.datasets.FashionMNIST(root='../datasets',  
                                                         train=True,  
                                                         download=True,  
                                                         transform=transform)  
  
fashion_mnist_dataloader = DataLoader(fashion_mnist_dataset, batch_size=8)
```



Source: https://www.tensorflow.org/datasets/catalog/fashion_mnist



Easy Network Creation

```
import torch.nn as nn
# defining the model
class Net(nn.Module):
    def __init__(self, input_size=1*28*28, output_size=100):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, output_size)

    def forward(self, x):
        x = self.fc1(x)
        return x

net = Net()
net = net.to(device)
```

Where is the
backward
pass?

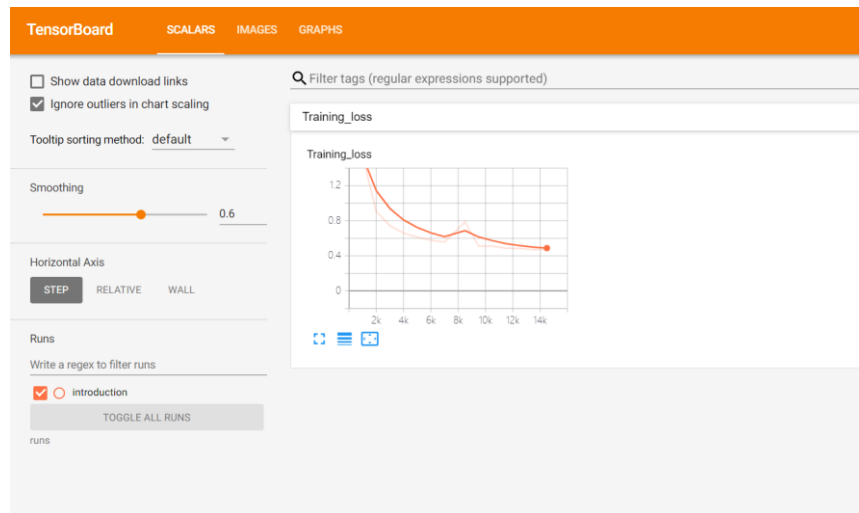


References on Pytorch

- Repository: <https://github.com/pytorch/pytorch>
- Examples (very nice):
<https://github.com/pytorch/examples>
- PyTorch for NumPy users:
<https://github.com/wkentaro/pytorch-for-numpy-users>

Tensorboard: Simple Logging

- Directly access tensorboard in your training loop
- Tensorboard generates the graph/timestamps etc. for you



TensorBoard: Visualize Networks

- Using a single forward pass, tensorflow can map and display your network graph

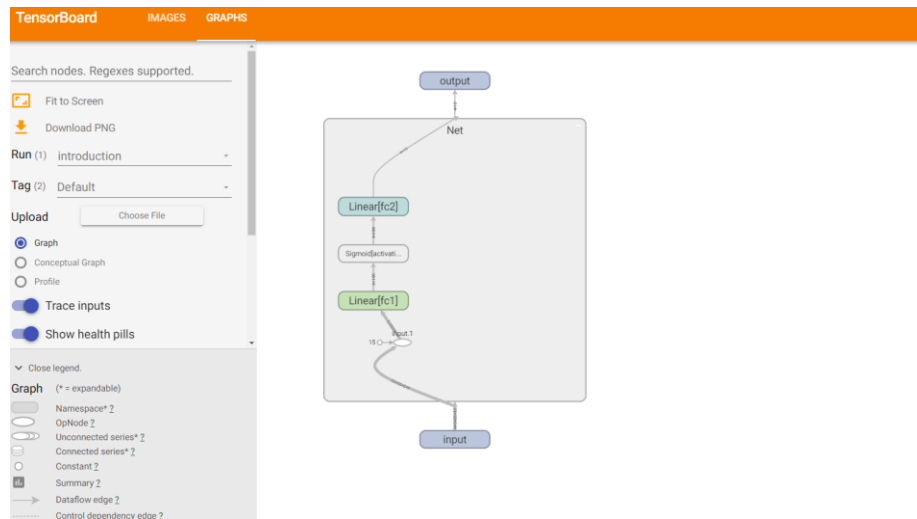
```
# Initialize your model
```

```
net = Net()
```

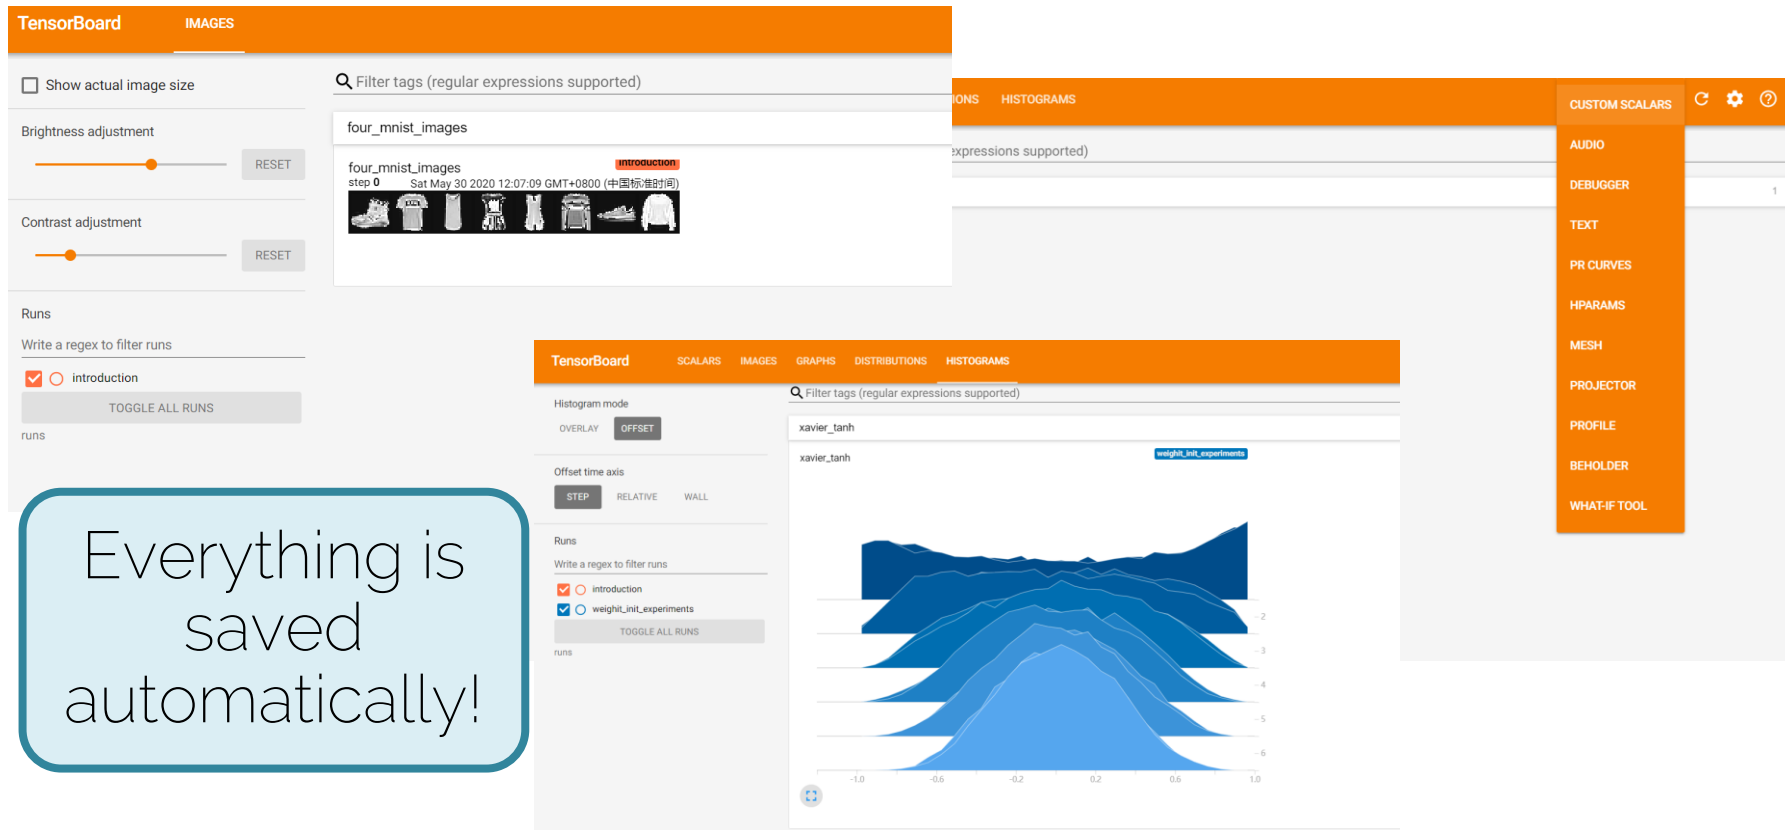
```
# Visualize its architecture in TensorBoard
```

```
writer.add_graph(net, images)
```

Graph creation
needs network
& one batch!

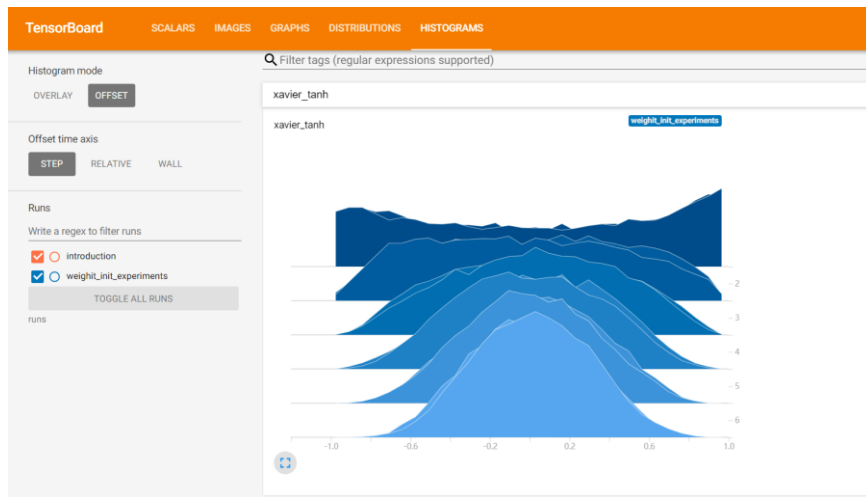



Document Everything!



Weight Initialization

- Histogram visualization for layer outputs can show off effects of weight initialization as shown in the lecture



Pytorch Lightning

Idea Behind PyTorch Lightning

Classify our code into three categories

1. **Research** code (the exciting part!, changes with new tasks, models etc.)

→ LightningModule

2. **Engineering** code (the same for all projects and models)

→ Trainer

3. **Non-essential** code (logging, organizing runs)

→ Callbacks

Lightning Module

PyTorch

```
# model
class Net(nn.Module):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

# train loader
mnist_train = MNIST(os.getcwd(), train=True, download=True,
                    transform=transforms.ToTensor())
mnist_train = DataLoader(mnist_train, batch_size=64)

net = Net()

# optimizer + scheduler
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3)
scheduler = Steplr(optimizer, step_size=1)

# train
for epoch in range(1, 100):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)

        loss.backward()
        optimizer.step()
    if batch_idx % args.log_interval == 0:
        print('Train Epoch: {} [{} / {}] ({:.0f}%) \t Loss: {:.6f}'.format(
            epoch, batch_idx * len(data), len(train_loader.dataset),
            100. * batch_idx / len(train_loader), loss.item()))
```

Methods that need to be implemented

- `__init__`
- `forward`
- `training_step`
- `configure_optimizers`

Optional methods

- `validation_step`
- `validation_end`
- `Data handling`

Trainer

PyTorch

PyTorch Lightning

```
# model
class Net(nn.Module):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

# train loader
mnist_train = MNIST(os.getcwd(), train=True, download=True,
                    transform=transforms.ToTensor())
mnist_train = DataLoader(mnist_train, batch_size=64)

net = Net()

# optimizer + scheduler
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3)
scheduler = StepLR(optimizer, step_size=1)

# train
for epoch in range(1, 100):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)

        loss.backward()
        optimizer.step()
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
```

```
# model
class Net(LightningModule):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

    def train_dataloader(self):
        mnist_train = MNIST(os.getcwd(), train=True, download=True,
                            transform=transforms.ToTensor())
        return DataLoader(mnist_train, batch_size=64)

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)
        scheduler = StepLR(optimizer, step_size=1)
        return optimizer, scheduler

    def training_step(self, batch, batch_idx):
        data, target = batch
        output = self.forward(data)
        loss = F.nll_loss(output, target)
        return loss

if __name__ == '__main__':
    net = Net()
    trainer = Trainer()
    trainer.fit(net)
```

1. Initialize the model with hyperparameters for training (e.g. as a dictionary)
2. Trainer contains all code relevant for training our neural networks
3. Call the method `.fit()` for training the network

That's all you need to train you model 😊

Pros/Cons

- **Advantages**
 - Better overview of the relevant code
 - Nice debugging features
 - Many automated options, like logging
- **Potential Problems**
 - Can have issues like any stock library
 - Not always straightforward to add features yourself

Your tasks

- Go over all notebooks
 - There is no submission or code to implement
- Experiment yourself!
- Check out documentations

PyTorch Governance	Docs > torch.nn	Docs > torch.nn
PyTorch Governance Persons of interest		
Python API		
torch		torch.nn
torch.nn		Parameters
torch.nn.functional		• Containers
torch.Tensor		• Convolution layers
Tensor Attributes		• Pooling layers
torch.autograd		• Padding layers
torch.cuda		• Non-linear activations (weighted sum, nonlinear)
torch.distributed		• Non-linear activations (other)
torch.distributions		• Normalization layers
torch.hub		• Recurrent layers
torch.jit		• Transformer layers
torch.nn.init		• Linear layers
torch.nnix		• Dropout layers
torch.optim		• Sparse layers
Quantization		• Distance functions
torch.random		• Loss functions
torch.sparse		• Vision layers
torch.storage		• DataParallel layers (multi-GPU, distributed)
torch.utils.bottleneck		• Utilities
torch.utils.checkpoint		• Quantized Functions
torch.utils.data		
torch.utils.dlpack		
torch.utils.model_zoo		


Optional Submission: CIFAR10 (again)

- We will open an optional leaderboard
 - TBD
- Rules (not enforced):
 - Use the previous Cifar10 dataset class in pytorch
 - No convolutional layers
- Use Piazza, especially if you have a successful approach, do a quick write up and discuss

Exam

- Exam date
 - 10.02.22, 10:00
- Exam Format
 - Online
- Mock Exam
 - Uploaded

Chair of Visual Computing & Artificial Intelligence
Department of Informatics
Technical University of Munich



Note:

- During the attendance check a sticker containing a unique code will be put on this exam.
- This code contains a unique number that associates this exam with your registration number.
- This number is printed both next to the code and to the signature field in the attendance check list.

Introduction to Deep Learning

Exam: IN2346 / endterm
Examiner: Prof. Dr. Matthias Nießner

Date: Tuesday 23rd February, 2021
Time: 08:00 – 09:30

	P 1	P 2	P 3	P 4	P 5	P 6	P 7
I							

Summary

- Monday 6.12: Watch Lecture 8
 - Training NNs 3
- No exercise submission
- Thursday 9.12: Tutorial 8
 - Autoencoder