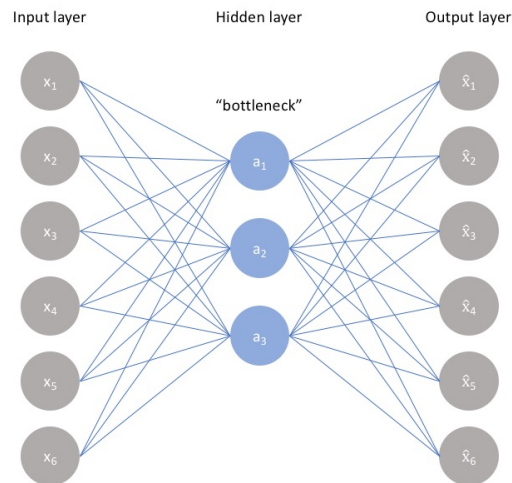


# Introduction to Deep Learning (I2DL)

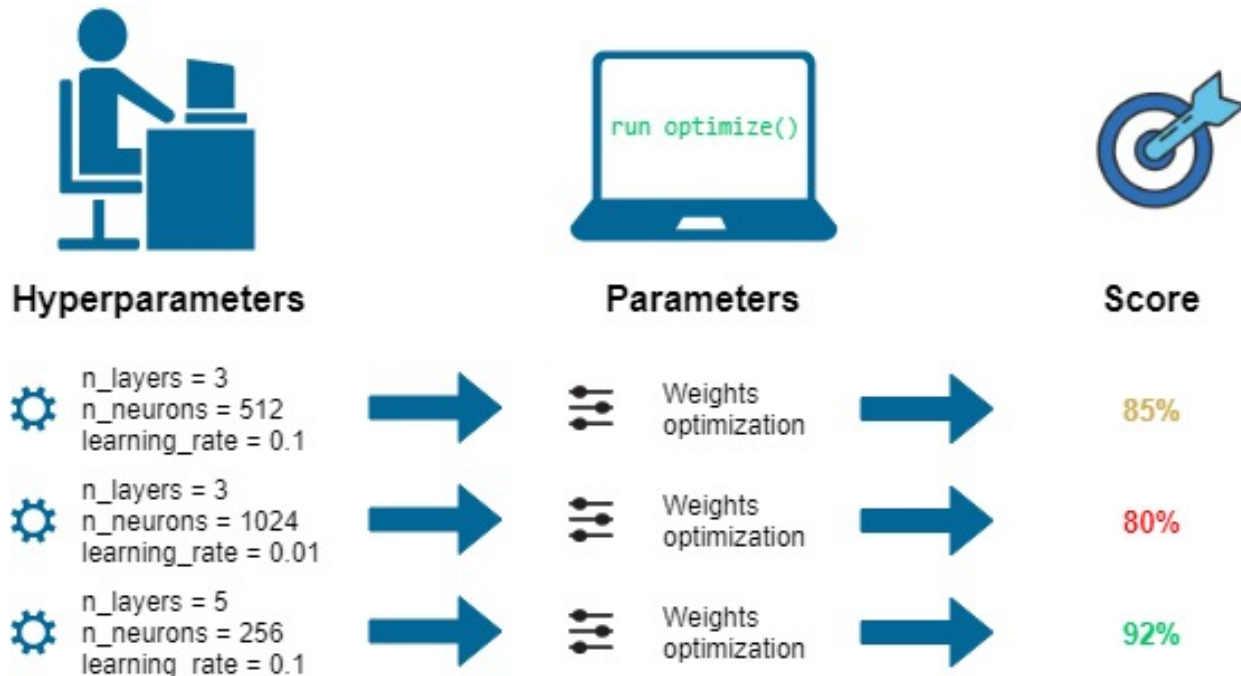
## Exercise 8: Autoencoder

# Today's Outline

- Hyperparameter tuning
- Exercise 8
  - Batch Normalization & Dropout
  - Transfer Learning
  - Autoencoder
- Personal: Github/Exposure



# Hyperparameter Tuning



Source: <https://images.deepai.org/glossary-terms/05c646fe1676490aa0b8cab0732a02b2/hyperparams.png>

# Hyperparameter Tuning

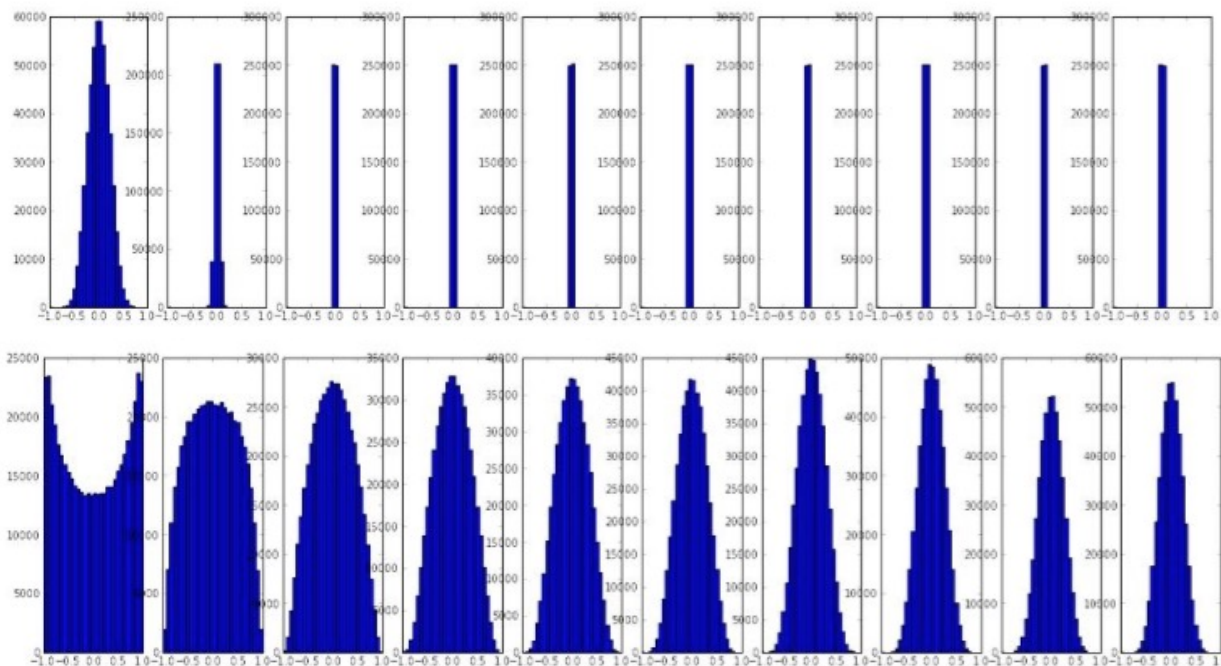
- Slides on Piazza
  - Check them out if you haven't done it yet
- It is important
  - Regardless of your resources
  - There is no all in one recipe
  - If you need more practice: optional submission on CIFAR10



# Improve your training!

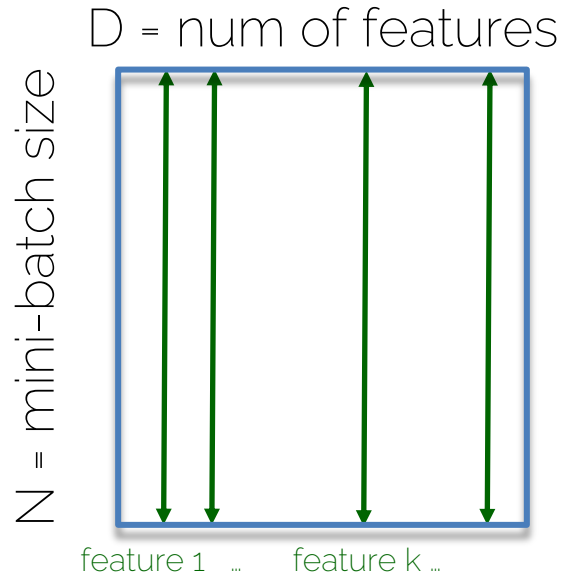
# Batch Normalization

- All we want is that our activations do not die out



# Batch Normalization

- Wish: Unit Gaussian activations



Mean of your mini-batch examples over feature k

Unit gaussian

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{\text{Var}[\mathbf{x}^{(k)}]}}$$

# Batch Normalization


- 1. Normalize

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

- 2. Allow the network to change the range

$$\mathbf{y}^{(k)} = \gamma^{(k)} \hat{\mathbf{x}}^{(k)} + \beta^{(k)}$$

backprop



The network can learn to undo the normalization

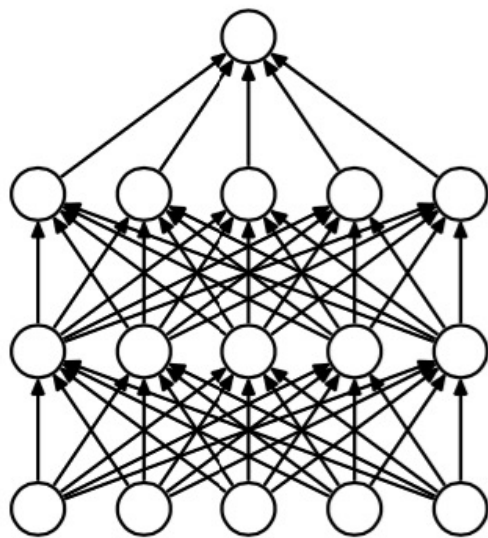
$$\gamma^{(k)} = \sqrt{Var[\mathbf{x}^{(k)}]}$$

$$\beta^{(k)} = E[\mathbf{x}^{(k)}]$$

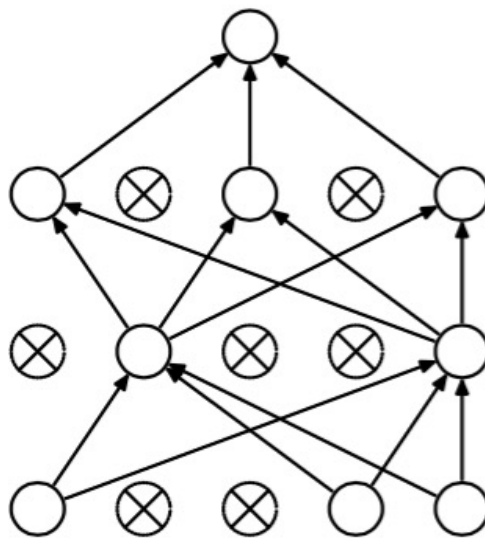


# Dropout

- Using half the network = half capacity



(a) Standard Neural Net

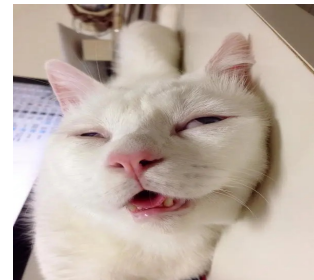
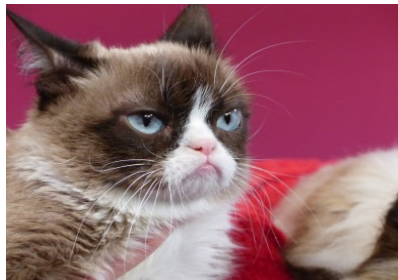


(b) After applying dropout.

Forward

# Transfer Learning

# Transfer Learning: Example Scenario

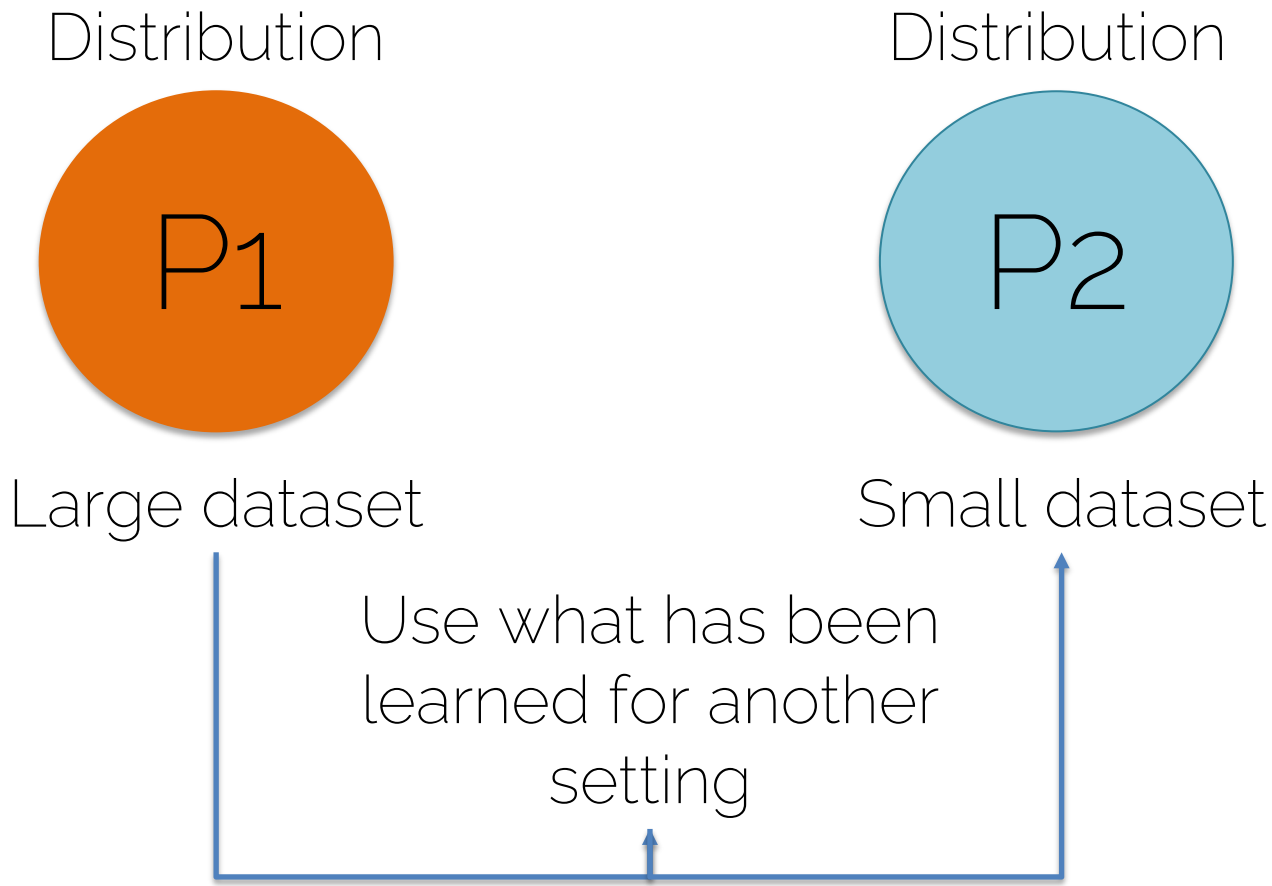


- Need to build a Cat classifier
- Only have a few images ~10 000

# Transfer Learning

- Problem Statement:
  - Training a Deep Neural Network needs a lot of data
  - Collecting much data is expensive or just not possible
- Idea:
  - Some problems/ tasks are closely related
  - Can we transfer knowledge from one task to another?
  - Can we re-use (at least parts of) a pre-trained network for the new task?

# Transfer Learning



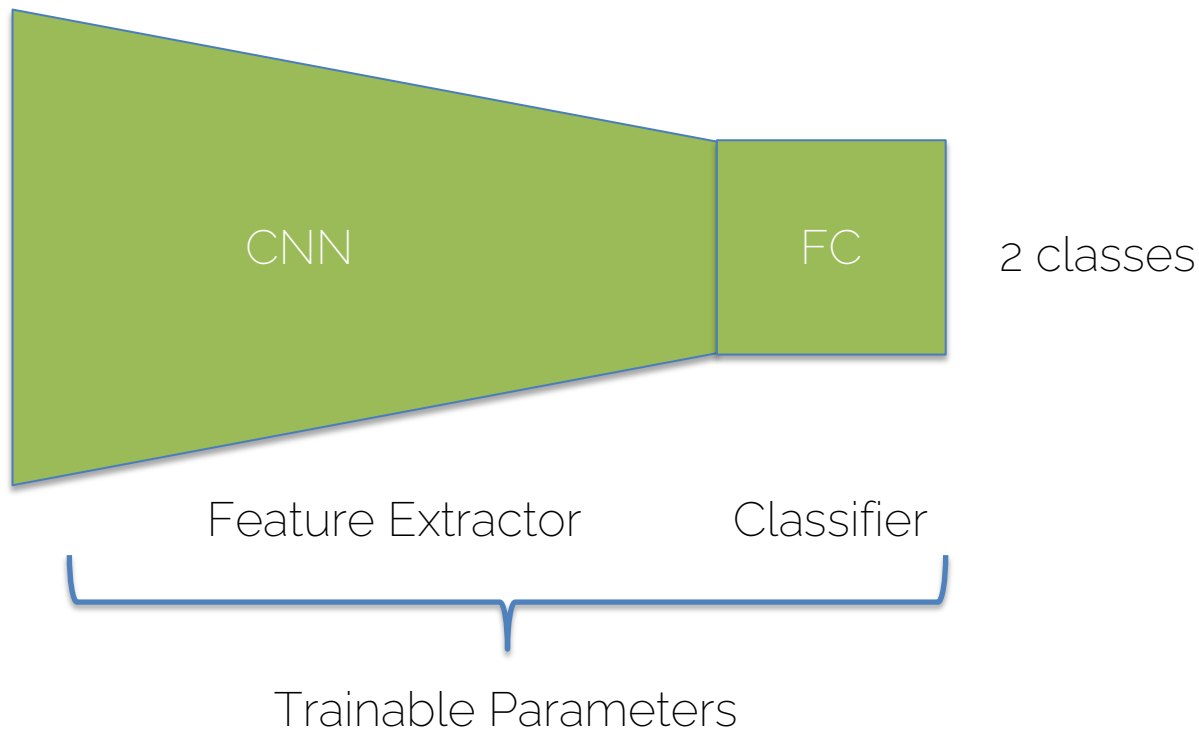
# Transfer Learning



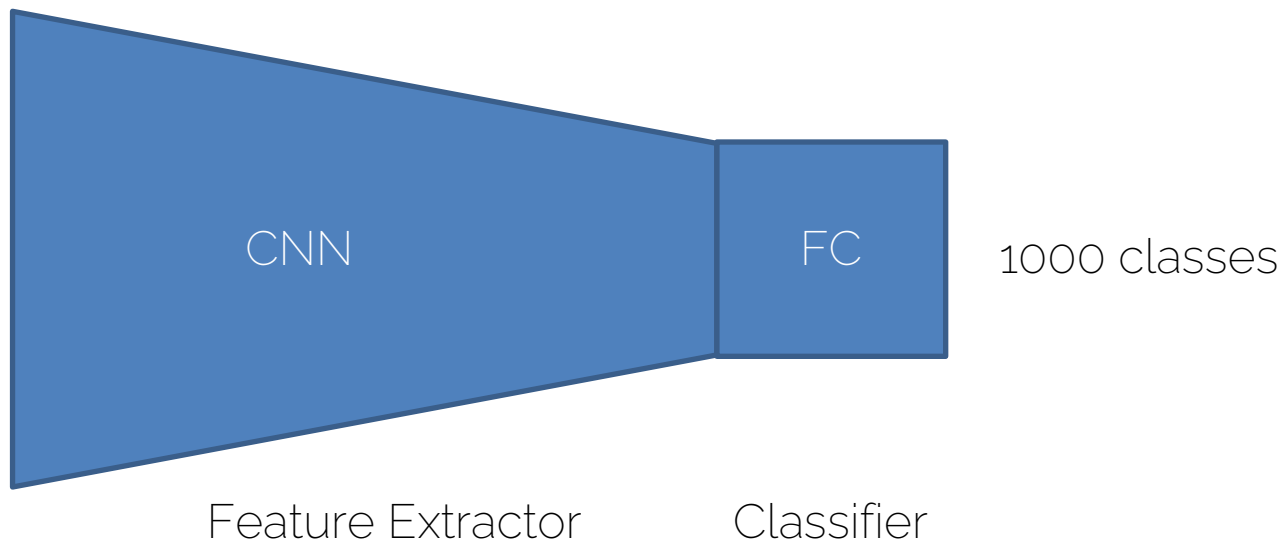
Coloring Legend:

 Untrained

 Trained



# Transfer Learning



Coloring Legend:

 Untrained

 Trained

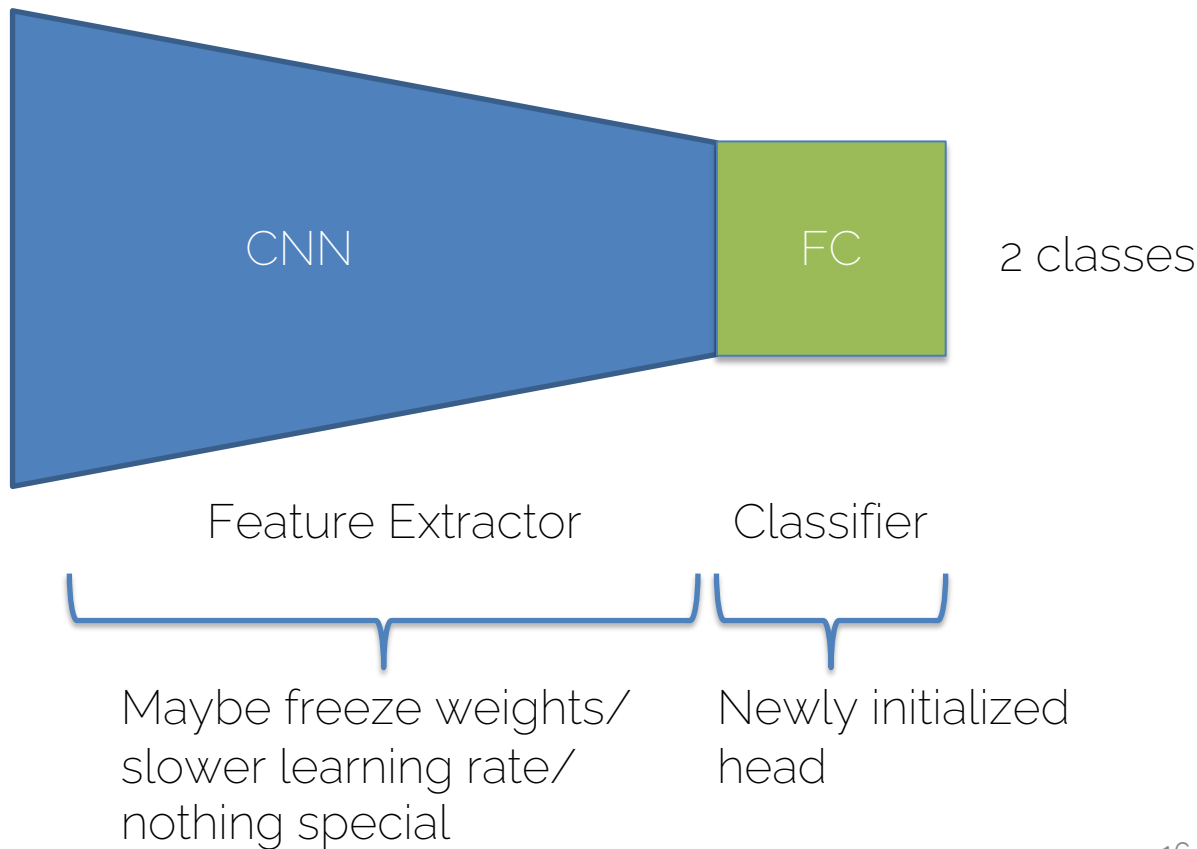
# Transfer Learning



Coloring Legend:

 Untrained

 Trained

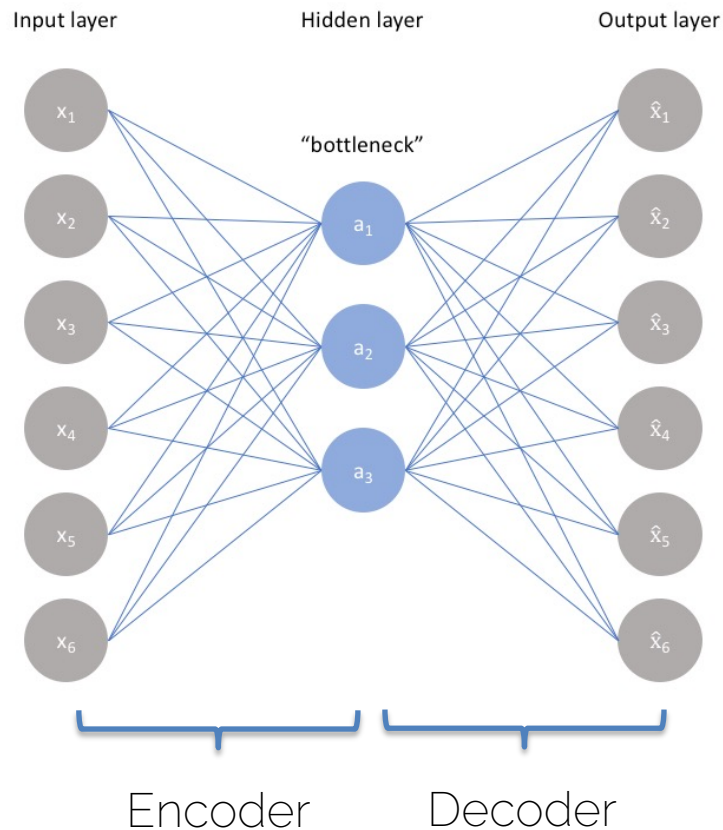




# Application: Autoencoder (Sub 8)

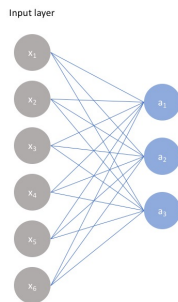
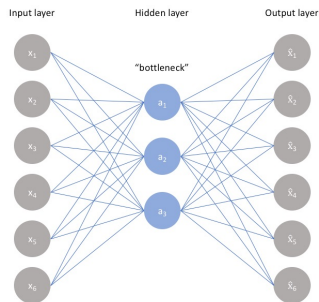
# Autoencoder

- Task
  - Reconstruct the input given a lower dimensional bottleneck
  - Loss: L1/L2 per pixel
- Actually need no labels!
- Without non-linearities: similar to PCA



# Transfer Using an Autoencoder

- Step 1:
  - Train an Autoencoder on a large (maybe unlabelled) dataset very similar to your target dataset
- Step 2:
  - Take pre-trained Autoencoder and use it as the first part of a classification architecture for your target dataset



# Personal Note: Github/Exposure

- Why do I want exposure or a portfolio?
- What is useful?
  - Something to talk about in interviews
- Posting I2DL solutions is not a helpful git for you
  - Maybe among other students...
- Projects:
  - internships/ guided research/ any task basically
  - Document your process (blog), show and visualize your data processing, discuss design decisions and publish code

# Summary

- Monday 13.12: Watch Lecture 9
  - CNN
- Exercise 8 Submission
  - Autoencoder: 15.12.2021 15.59
- Thursday Tutorial 9: 16.12.2021
  - Facial Keypoint Detection

See you next week!