

Introduction to Deep Learning (I2DL)

Exercise 11: Sentiment Analysis with RNNs

Today's Outline

- Online Exam: February 10, 2021 (register by 15.1)
- Exercise 10: Segmentation
Case Study of Student's Solutions
- RNNs and LSTMs
- Exercise 11: Sentiment Analysis with RNNs

Exam

- Due to current university restrictions: Online exam
 - No retake exam
 - Bonus will be transferred to any future version of this class
- **Organization:** We will release more detailed information on Piazza as soon as possible
- **Exam relevant content:**
 - Lectures
 - Exercises including optional notebooks

Case Study:

Exercise 10 Semantic Segmentation

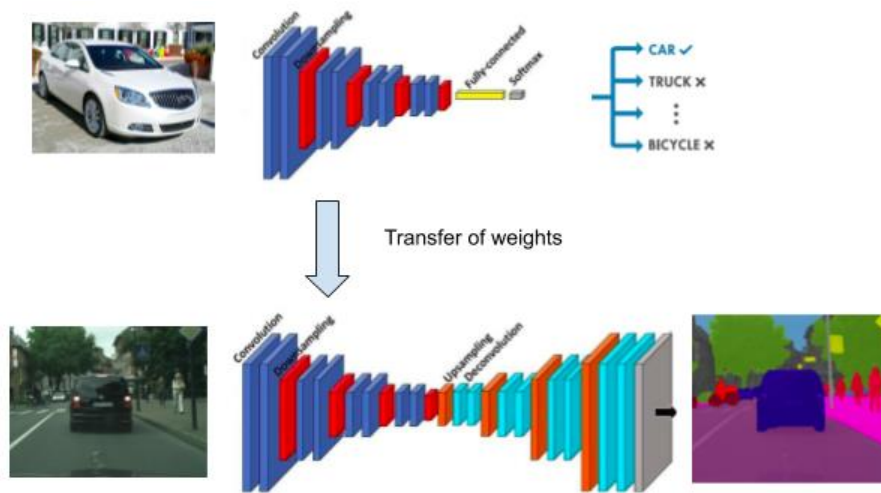
Exercise 10: Semantic Segmentation

- **Goal:** Assign a label to each pixel of the image
- **Output of the network:** Segmentation mask with same shape as input image
- **Dataset:** MSRC v2 dataset, 23 object classes, contains 591 images with accurate pixel-wise labeled images



Transfer Learning for Segmentation

- Idea: Encoder-Decoder Architecture
- **Transfer Learning:** CNNs trained for image classification contain meaningful information that can be used for segmentation -> Encoder
- **Check out:** pre-trained networks like AlexNet, MobileNets



Leaderboard Submission #10

WS 20/21

Rank	User	Score	Pass
#1	w1459	91.49	✓
#2	w1691	90.80	✓
#3	w1492	90.65	✓
#4	w1258	89.12	✓
#5	w1631	89.06	✓
#6	w1594	88.88	✓
#7	w1463	87.93	✓
#8	w1284	87.46	✓
#9	w1679	87.19	✓
#10	w1468	86.71	✓

SS 21

#	User	Score
1	u0586	92.06
2	u1133	91.72
3	u1055	91.38
4	u0330	91.35
5	u0642	90.87
6	u0623	90.62
7	u0916	89.67
8	u0263	89.31
9	u0255	89.10
10	u0048	89.09

23/06, 20:00

Top 1 Submission: 92,06%

```
self.model_ft = torchvision.models.mobilenet_v2(pretrained=True).features
num_fts = 1280 # 1280 of 8x8 features

self.decoder = nn.Sequential(
    nn.Upsample(scale_factor=2, mode='bicubic', align_corners=True),
    nn.ConvTranspose2d(1280, 192, kernel_size=2, stride=2),
    nn.Conv2d(192, 192, kernel_size=3),
    nn.BatchNorm2d(192),
    nn.ReLU(),

    nn.Upsample(scale_factor=2, mode='bicubic', align_corners=True),
    nn.ConvTranspose2d(192, 80, kernel_size=2, stride=2),
    nn.Conv2d(80, 80, kernel_size=3, padding=1),
    nn.BatchNorm2d(80),
    nn.ReLU(),

    nn.Upsample(scale_factor=2, mode='bicubic', align_corners=True),
    nn.Conv2d(80, num_classes, kernel_size=1),
)
```

Summary:

- Architecture: Encoder-Decoder
- Encoder: Pre-trained Mobilenet
- Decoder: Several Blocks consisting of
 - Upsampling Layers
 - Transposed Conv
 - BatchNorm
 - ReLU

Top 2 Submission: 91, 72%

```
self.encoder = models.mobilenet.mobilenet_v2(pretrained=True).features
```

```
self.inner = nn.Sequential(  
    nn.Conv2d(1280, l1_depth, 1, 1, 0),  
    nn.BatchNorm2d(l1_depth),  
    nn.ReLU(inplace=True),  
    nn.Conv2d(l1_depth, l1_depth, 3, 1, 1),  
    nn.BatchNorm2d(l1_depth),  
    nn.ReLU(inplace=True)  
)
```

```
self.dec_upsample_layers = nn.Sequential(  
    nn.ConvTranspose2d(l1_depth, l2_depth, 3, 2, 1),#1-2  
    nn.ConvTranspose2d(l2_depth, l3_depth, 3, 2, 1),#2-3  
    nn.ConvTranspose2d(l3_depth, l4_depth, 3, 2, 1),#3-4  
    nn.ConvTranspose2d(l4_depth, l5_depth, 3, 2, 1),#4-5  
    nn.ConvTranspose2d(l5_depth, l6_depth, 3, 2, 1),#5-6  
)
```

```
self.dec_layers = nn.Sequential(  
    nn.Sequential(#12  
        nn.ReLU(),  
        #nn.Dropout2d(),  
        nn.Conv2d(l2_depth+96, l2_depth, 3, 1, 1),#input: 100x15x15  
        nn.BatchNorm2d(l2_depth),  
        nn.ReLU(),  
        #nn.Dropout2d(),  
  
        nn.Conv2d(l2_depth, l2_depth, 3, 1, 1),#input: 100x15x15  
        nn.BatchNorm2d(l2_depth),  
        nn.ReLU(),  
        #nn.Dropout2d(),  
    ),
```

```
    nn.Sequential(#13  
        nn.ReLU(),  
        #nn.Dropout2d(),  
        nn.Conv2d(l3_depth+32, l3_depth, 3, 1, 1),#input: 80x30x30  
        nn.BatchNorm2d(l3_depth),  
        nn.ReLU(),  
        #nn.Dropout2d(),  
  
        nn.Conv2d(l3_depth, l3_depth, 3, 1, 1),#input: 80x30x30  
        nn.BatchNorm2d(l3_depth),  
        nn.ReLU(),  
        #nn.Dropout2d(),  
    ),
```

...

Top 2 Submission: 91, 72%

```
def forward(self, x):
    #####
    #                               YOUR CODE                               #
    #####

    enc_layer_res = []

    res = x
    for i, layer in enumerate(self.encoder):
        new_res = layer(res)

        if res.shape[-2:] != new_res.shape[-2:]:
            enc_layer_res.append(new_res)
        else:
            enc_layer_res[-1] = new_res

        res = new_res

    res = self.inner(res)

    for i, layer in enumerate(self.dec_upsample_layers):
        if len(enc_layer_res)-1 == i:
            output_size = (240, 240)
        else:
            output_size = enc_layer_res[-i-2].shape[-2:]


        res = layer(res, output_size=output_size)

        if len(enc_layer_res)-1 != i:
            res = torch.cat([res, enc_layer_res[-i-2]], dim=1)

        res = self.dec_layers[i](res)

    x = res

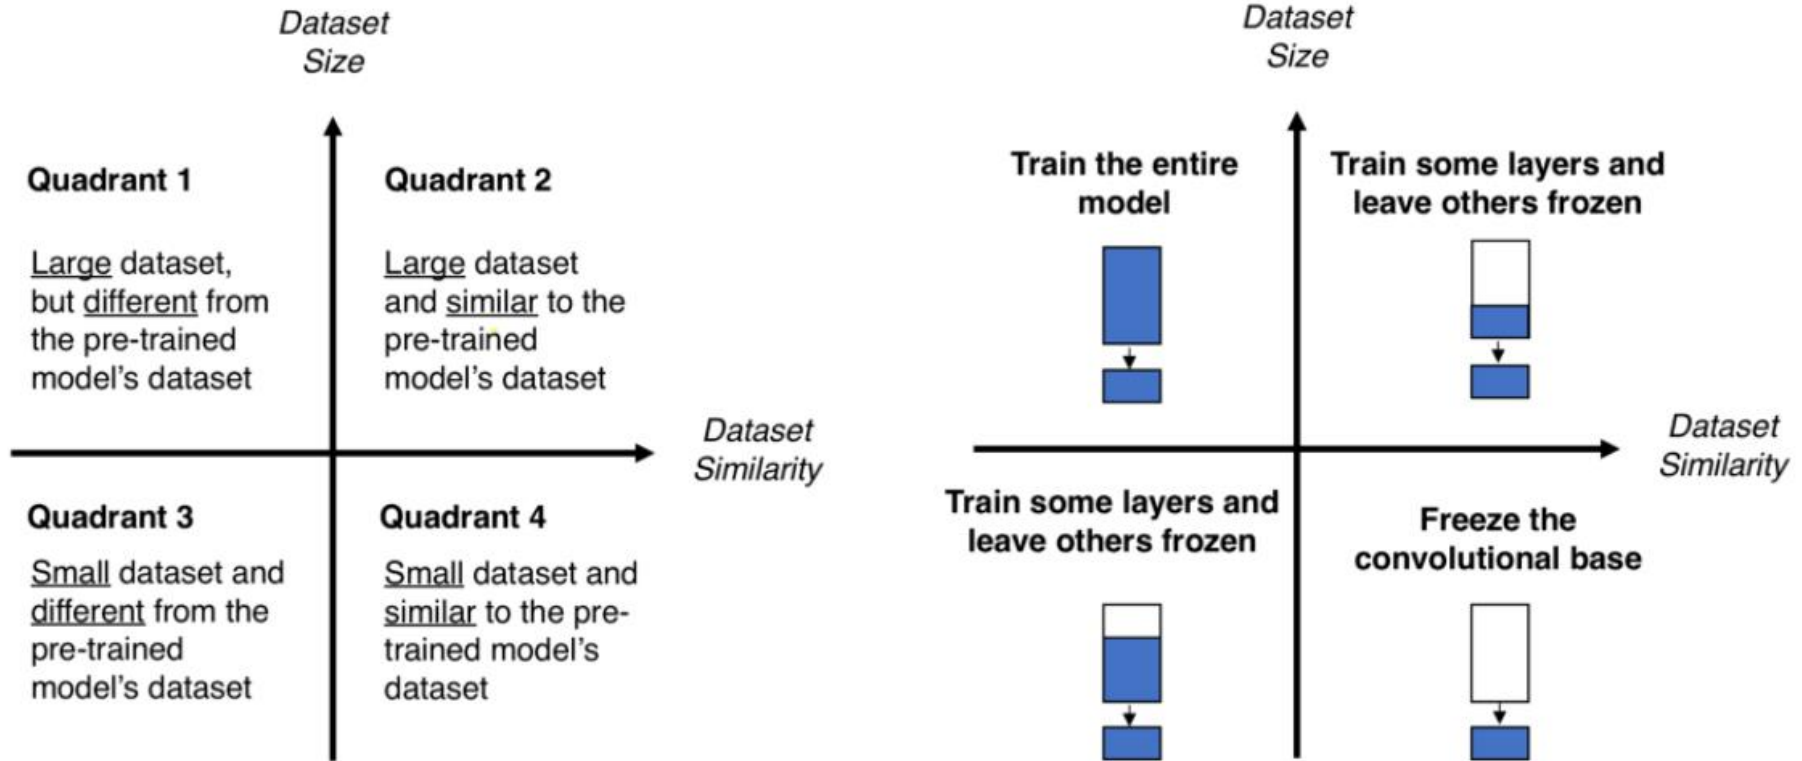
    #####
    #                               END OF YOUR CODE                               #
    #####
```



Summary:

- Architecture: UNet
- Forward pass:
 - Concatenate encoder features in decoder

When/What to Fine-tune.



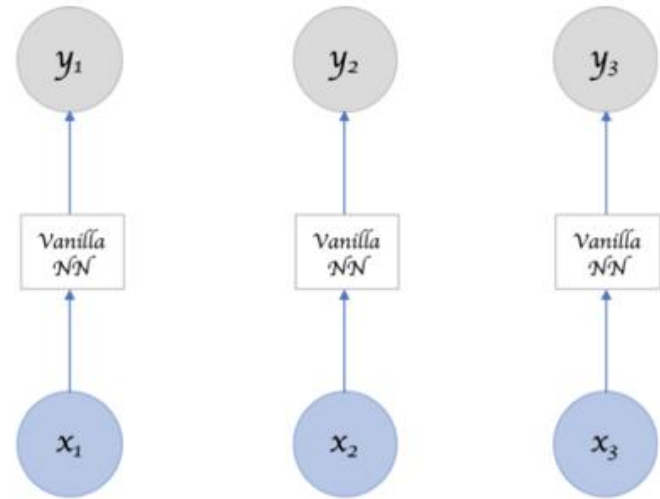
Recurrent Neural Networks

Vanilla Feed-Forward Neural Networks

- **Input:** Fixed size vector as input, e.g. images
- **New type of input:** Series of input with no predetermined limit on size
- **Examples:** Series of images = Video, Series of words = Text

IDEA: Use vanilla network repeatedly for a series of input?

PROBLEM: Series of inputs mean that the components are related to each other




Recurrent Neural Networks

- **Idea:** Network that can capture the relationship between the inputs
- **RNNs:** Learning process is not independent
 - Remember things from processing trainings data
 - Remember things learnt from prior inputs, prior inputs influences decision
- **In other words:** RNNs produce different outputs for same input depending on previous outputs in the series.



$$\mathbf{A}_t = \theta_c \mathbf{A}_{t-1} + \theta_x \mathbf{x}_t$$

Previous hidden state input



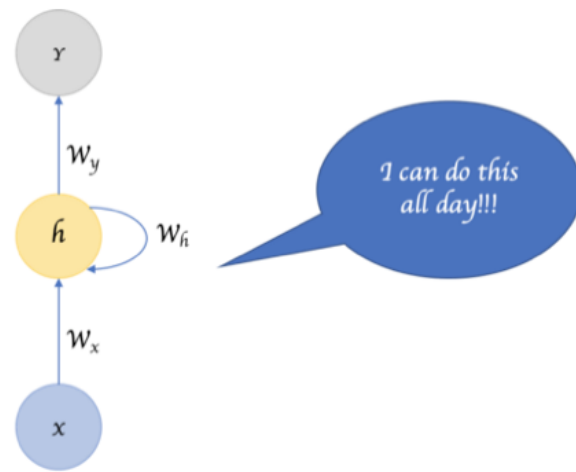
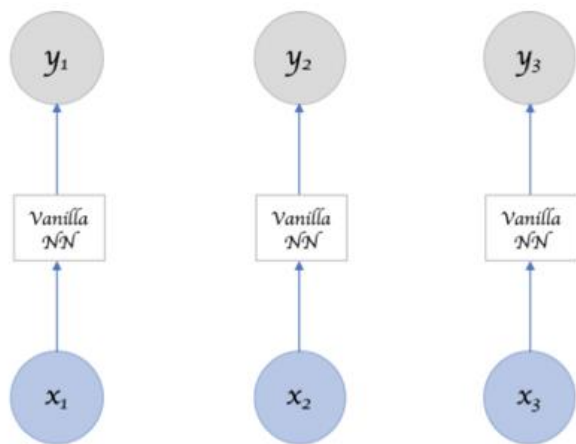
RNNs: Parameter sharing

- **Parameter Sharing:** Weights are shared across all inputs
- **Advantage:** Using same weights allows to process series of input without predefined length



RNNs: Parameter sharing

- **Parameter Sharing:** Weights are shared across all inputs
- **Advantage:** Using same weights allows to process series of input without predefined length



RNNs: Parameter sharing

- **Parameter Sharing:** Weights are shared across all inputs
- **Advantage:** Using same weights allows to process series of input without predefined length
- **Hidden state:** Ensures that each input undergoes a different procedure even with shared weights



RNN Concepts

one to one




Image Classification

one to many






Image Captioning
(image \rightarrow seq of words)

many to one




Sentiment Analysis
(seq of words \rightarrow sentiment)

many to many



Machine Translation (seq of words \rightarrow seq of words)


many to many




Video Classification on frame level (seq of frames \rightarrow seq of class.)

RNNs and LSTMs

- Disadvantage: RNNs struggle to learn long-term dependencies
- LSTMs (Long Short Term Memory Networks): special kind of RNNs capable of learning long-term dependencies



The *clouds* are in the *sky*.



I moved to *Germany* ... so I speak *German* fluently.

Exercise 11: Sentiment Analysis

Exercise 11: Goal

Review: I wouldn't rent this one even on dollar rental night.

Sentiment:



Review: Adrian Pasdar is excellent in this film. He makes a fascinating woman.

Sentiment:



Exercise 11: Content

- Optional Notebook: RNNs and LSTMs
- Notebook 1: Text Preprocessing and Embedding
- Notebook 2: Sentiment Analysis

Review: I wouldn't rent this one even on dollar rental night.

Sentiment:



Review: Adrian Pasdar is excellent in this film. He makes a fascinating woman.

Sentiment:



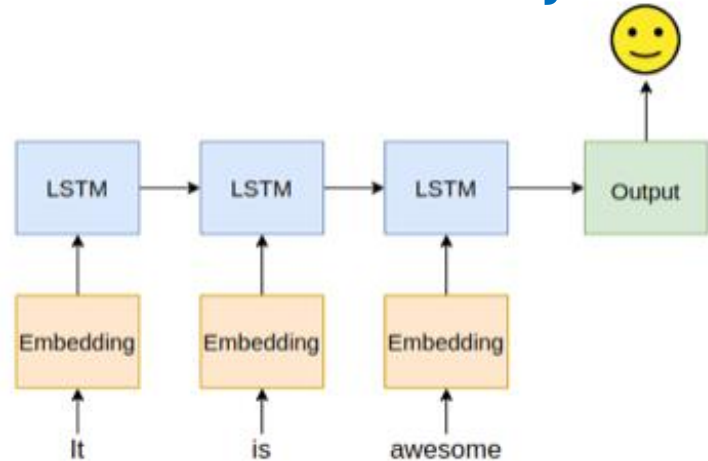
Notebook 1: Text Preprocessing and Embedding

- Sequential Data: from image data to text data
- Dataset: IMDb sentiment analysis dataset
- Goal of the notebook:
 - Data preparation
 - Implementation of Embedding layer



Notebook 2: Sentiment Analysis

- Network Architecture:
 - Embedding layer
 - RNN
 - Output layer, e.g. fully-connected layer
- Loss: Cross-Entropy Loss
- Performance measure: Accuracy
- Goal of the notebook: Implement and train a recurrent neural network for sentiment analysis



Submission Details

- Submission **Start**: January 13, 2021 13.00
- Submission **Deadline** : January 19, 2021 18.59
- Submit your trained model
- Your model's **accuracy on our test set** is all that counts in order to pass this submission!
 - Threshold to pass: 83%

Summary

- **Monday 17.1:** Watch Lecture 12
 - Advanced DL Topics
- **Wednesday 19.1:** Submit exercise
- This is the last tutorial

Good luck with the
exercise! 😊